How Governance Level Shapes AI Risk Mitigation Effectiveness: A Cross-Sector Comparative Analysis of AI Policy Frameworks

` As generative AI expands into all facets of education, government, and industry, each sector stewards policy development in unique manners. What began as experimental code in research labs now drafts legal briefs, grades essays, recommends medical treatments, and negotiates financial trades. With systems advancing faster than the statutes and syllabi meant to guide them, AI governance has shifted from a theoretical debate to an urgent question of who sets the rules, how those rules are enforced, and whether they meaningfully contain risk before harm occurs. This analysis examines how governance authority level—voluntary corporate self-regulation, institutional academic policy, and state legal mandate—shapes the enforceability, transparency, and effectiveness of AI risk mitigation.

To ground this comparison, the paper analyzes three distinct governance models. Anthropic's Responsible Scaling Policy (RSP v2.2, May 2025) represents a voluntary industry framework designed to manage risks from frontier AI model development through internal evaluation thresholds, staged deployment controls, and public commitments to safety research. Ohio State University's AI Use Guidance reflects an institutional education policy that governs how students may use AI tools in coursework, balancing academic integrity, pedagogical flexibility, and faculty discretion. By contrast, Texas HB 149—the Texas Responsible AI Governance Act, effective January 2026—embodies a state-level consumer protection statute that imposes legally binding requirements on AI developers and deployers operating within Texas, including compliance obligations and potential penalties.

At the heart of this comparison lies a central tension: whether aspirational principles backed by voluntary compliance can adequately address AI risks, or whether enforceable operational standards with legal or institutional authority are necessary to protect stakeholders. Corporate frameworks may adapt rapidly but lack external compulsion. Universities may embed norms within communities yet depend on decentralized enforcement. State law carries coercive power, but often at the cost of agility. The question is not merely which model exists, but which model meaningfully constrains risk.

This paper proceeds by first examining how each policy defines AI risk and allocates responsibility, revealing how authority level influences the scope of oversight. It then evaluates enforcement mechanisms and transparency requirements, comparing internal auditing, academic honor systems, and statutory penalties. Finally, it assesses practical effectiveness by analyzing implementation feasibility, compliance incentives, and stakeholder protection outcomes, drawing broader implications for cross-sector AI governance design.

Why this matters extends beyond theory. Professionals deploying AI systems must navigate overlapping corporate guidelines and state regulations. Policymakers face pressure to legislate without stifling innovation. Educators must design curricula that prepare students for AI-integrated workplaces while preserving academic integrity. By comparing governance across industry, academia, and state law, this analysis clarifies how authority structure shapes real-world accountability in an era where algorithmic decisions increasingly shape human opportunity.

Before comparing what these frameworks do well or poorly, it helps to get grounded in what each one actually is, who made it, who it applies to, and what it's trying to govern. These three policies look very different on the surface, and that's partly because they come from completely different worlds.

Anthropic is an AI safety company, and its Responsible Scaling Policy (RSP, currently version 2.2 as of May 2025) is a document the company wrote and imposed on itself. There's no outside agency requiring this, it's a voluntary commitment backed by Anthropic's Board of Directors and overseen by its Long-Term Benefit Trust. Because it's self-imposed, the only people who have to follow it are Anthropic employees and the teams building its models. No external developer, partner, or end user is bound by it.

In terms of scope, the RSP focuses on what Anthropic calls "frontier AI", the most powerful models being developed or trained at any given time, particularly those with potential to cause large-scale harm in areas like bioweapons (CBRN threats) or autonomous AI systems capable of operating without human control. The core mechanism is a tiered system of AI Safety Levels (ASL-1 through ASL-4), where each level corresponds to a set of capability thresholds. Once a model hits a threshold, specific safeguards kick in before development can continue. As the RSP puts it, its goal is "not to train or deploy models capable of causing catastrophic harm unless we have implemented safety and security measures." The policy is designed to be preventive, catching risk before deployment rather than reacting after something goes wrong.

OSU's guidance comes from a very different place. It was developed by the university's Office of Distance Education as a resource for instructors trying to figure out what to do about students using tools like ChatGPT. The document doesn't carry the weight of a standalone law, its authority flows from OSU's existing Code of Academic Misconduct and the broader institutional commitment to academic integrity. In practice, that means instructors are expected to set AI policies at the assignment level, and students are expected to follow them.

The scope here is tied specifically to academic coursework: things like grammar checkers, text generators, summarizers, and image creation tools used by students on assignments. The guiding mechanism is instructor discretion; there's no single university-wide rule about what AI is or isn't allowed, because the guidance recognizes that what makes sense in a creative writing class is different from what makes sense in a statistics course. When AI use is permitted, students are required to cite the tool, explain how they used it, and describe its influence on the final product. The stated purpose is to maintain academic integrity while recognizing AI tools' pedagogical potential, and to address equity concerns for students who may not have equal access to these tools.

Texas HB 149, the Texas Responsible AI Governance Act, is in a different category altogether. This is state law, passed by the Texas Legislature and effective January 2026. It's not a suggestion or a set of guidelines; it's a legal mandate with civil penalties attached. Anyone conducting business in Texas, developing or deploying AI systems used by Texas residents, or operating as a state government agency falls within its reach.

The scope is broader than either of the other two frameworks: it covers consumer-facing AI systems, government deployment of AI, high-risk uses like biometric surveillance and social scoring, and explicitly prohibited behaviors like using AI to psychologically manipulate users or discriminate against protected classes. The mechanism combines prohibition with mandatory disclosure; covered entities must clearly inform users when they're interacting with AI, before or at the moment of interaction. Enforcement sits with the Texas Attorney General, who can investigate complaints, assess civil penalties, and refer cases for further action. A regulatory sandbox provision also allows for controlled testing of new AI applications under modified requirements. The bill's stated goal is to "protect public safety, individual rights, and privacy while encouraging the safe advancement of AI technology."

What the table below makes clear is that these three frameworks don't just differ in what they cover, they differ in who has the power to make anything happen. That distinction is going to matter a lot in the sections that follow.

Table 1: Scope, Enforcement, and Stakeholders Across Regulatory Tiers

	Anthropic RSP	**OSU Policy**	**Texas HB 149**
Authority Type	Voluntary	Institutional	Legal Mandate
Enforcement Entity	Internal (RSO, Board)	Faculty/COAM	Attorney General
Primary Stakeholder	AI Developers	Students/Instructors	Consumers/Public
Risk Focus	Catastrophic/Existential	Academic Integrity	Consumer Protection

Definitions shape what gets regulated. Anthropic's technical precision targets capability-level harms but misses structural bias embedded in everyday applications. OSU's broad equity framing raises the right concerns but provides no enforcement mechanism. Texas's legal intent standard creates formal anti-discrimination protection while simultaneously making it nearly unenforceable against algorithmic systems whose discriminatory outputs emerge from training data rather than deliberate design. No single framework captures the full scope of

AI-related fairness harm—each governs only the dimension of bias legible within its sector's existing vocabulary.

Transparency serves three distinct governance functions across these frameworks—corporate legitimacy, academic integrity, and consumer protection—and the mechanism each policy deploys to achieve it reveals how much disclosure is actually required versus merely encouraged. Anthropic's RSP operationalizes transparency as a voluntary, self-directed commitment: Section 7.2 commits to publicly releasing capability and safeguards reports when models are deployed, soliciting external expert input, and notifying U.S. government entities when models surpass ASL-2 thresholds—yet sensitive information is redacted at Anthropic's discretion, external consultation remains optional, and no external party can compel disclosure or impose penalties for its absence. This is transparency designed to build institutional credibility and model industry norms, not to satisfy an enforceable public right to information. OSU's framework shifts transparency inward toward the academic community: when AI use is permitted, students must cite the specific tool used, its purpose, and the extent of its influence on the final product—and instructors are expected to communicate the rationale behind their AI permissions at the assignment level. This creates meaningful disclosure obligations within the classroom, but says nothing about how OSU itself, or the third-party AI platforms students are directed to use, handle the data generated through those interactions. Texas HB 149 represents the only framework where transparency is externally enforceable against covered entities: Section 552.051 mandates that any governmental agency deploying a consumer-facing AI system must disclose this fact clearly, conspicuously, and in plain language—without dark patterns—"before or at the time of interaction," with violations subject to civil penalties reaching $200,000 and a public complaint mechanism administered by the

Attorney General (Section 552.102). Taken together, these three approaches reveal transparency requirements vary so dramatically across governance contexts that they produce incompatible disclosure regimes—voluntary commitments bend to corporate willingness, institutional policies protect students within bounded academic communities, and only legal mandates create durable, externally enforceable accountability to the public.

Privacy emerges across these three frameworks not as a unified principle but as a sectoral preoccupation—each policy protects the data most politically consequential to its governing context while leaving adjacent privacy concerns largely unaddressed. Anthropic's RSP constructs its most sophisticated privacy architecture around model weights rather than user data: the ASL-2 Security Standard mandates encryption, compartmentalization, and access controls, while the ASL-3 Security Standard escalates these protections against non-state attackers attempting model theft (RSP Section 4.2, Appendix B). This architecture reflects an industry priority—proprietary model parameters represent enormous competitive and safety value—but the RSP does not substantively govern what happens to the data users generate through model interactions, a gap that affects millions of consumer touchpoints daily. OSU's guidance acknowledges the privacy risks students assume when submitting work to third-party AI platforms, directing instructors to OSU's Security and Privacy Statement on Artificial Intelligence, but it creates no new privacy requirements and relies entirely on existing institutional protections like FERPA—frameworks designed for human-generated student records, not AI-processed behavioral data. Texas HB 149 addresses privacy most directly at the consumer interface: Section 552.054 prohibits government entities from deploying AI to uniquely identify individuals through biometric data without consent, and Section 552.053 bans government social scoring systems that aggregate behavioral data to produce consequential classifications. Together, these provisions target the privacy harms most visible in public-sector AI deployment, but they do not extend equivalent protections to private-sector data collection practices beyond what existing consumer data law (Chapter 541) already covers. The pattern across all three frameworks confirms that privacy protection follows sectoral salience—model weight integrity for industry, student record stewardship for education, and biometric identity for government—producing a fragmented landscape where the data each sector ignores is precisely what another leaves ungoverned.

Where privacy protections reveal what each sector values most, accountability structures reveal who holds the power to enforce those values—and crucially, who is excluded from doing so. Anthropic's RSP constructs the most architecturally elaborate accountability system of the three frameworks, layering the Responsible Scaling Officer's operational authority beneath Board of Directors oversight and Long-Term Benefit Trust consultation, while supplementing these with anonymous employee reporting channels and annual third-party procedural compliance reviews (RSP Section 7.1). This structure is sophisticated but entirely self-contained: the RSO reports upward to the Board, external experts provide input at Anthropic's discretion, and no consumer, regulator, or affected community holds standing to initiate an accountability process or challenge a deployment decision. OSU's accountability architecture distributes authority differently but remains similarly bounded—instructors hold primary sanctioning power over AI misuse within their courses, and the Committee on Academic Misconduct adjudicates escalated integrity violations through an established institutional process. Student complaint channels exist, but accountability flows through the academic hierarchy rather than outward, meaning non-students and the broader public have no mechanism to contest how OSU's institutional AI policies operate. Texas HB 149 breaks this pattern entirely by vesting exclusive enforcement authority in the Attorney General, empowering consumers to submit complaints through a publicly maintained online mechanism (Section 552.102), authorizing civil penalties ranging from $10,000 to $200,000 per violation (Section 552.105), and enabling state agency sanctions including license revocation for licensed entities (Section 552.106)—though the 60-day cure period (Section 552.104) means documented harms may continue without penalty while violators remediate. Anthropic and OSU both construct real, layered accountability structures—but ones that terminate within their own institutions, leaving no mechanism for consumers or affected communities to initiate review. Only Texas creates accountability that travels outward: Attorney General enforcement, civil penalties up to $200,000, and a public complaint mechanism mean that the person harmed, not the institution that caused the harm, can trigger consequences—making the difference between these frameworks not merely procedural, but a question of whose experience of AI harm counts.

Human oversight is operationalized most differently across these three frameworks than any other principle, exposing a fundamental tension between governance ideals and the automation trajectories each sector is simultaneously pursuing. Anthropic's RSP embeds human review at multiple deployment decision points—red teaming by external experts stress-tests model capabilities before release, capability assessments require CEO and Responsible Scaling Officer approval, and Board of Directors notification is mandatory before scaling beyond established thresholds (RSP Sections 3.3, 4.3)—yet the same policy that mandates this oversight is actively developing toward the AI R&D-4 capability threshold, defined as the ability to fully automate the work of an entry-level Anthropic researcher, a development that would structurally erode the human judgment these procedures currently depend on. OSU's guidance operationalizes human oversight most robustly and without this internal contradiction: instructor authority over AI use permissions is absolute and non-delegable, the guidance consistently frames human engagement as pedagogically necessary rather than merely procedurally required, and the emphasis on authentic skill development treats human judgment as the irreducible goal of education rather than a checkpoint within it. Texas HB 149 establishes human oversight through prohibition rather than process—Section 552.052 bars AI systems designed to manipulate users toward self-harm or criminal behavior, drawing a legal floor beneath which automated influence cannot descend—but stops short of mandating affirmative human review requirements for consequential AI-assisted decisions. The comparative insight is sharp: human oversight is most structurally secure where automation pressures are lowest, most aspirationally constructed where they are highest, and most legally enforced where consumer vulnerability is the governing concern.

Table 2: Ethical Principle Operationalization Across AI Governance Frameworks

Ethical Principle	**Anthropic RSP**	**OSU Policy**	**Texas HB 149**
Fairness & Bias Mitigation	✓	✓	✓
Transparency	✓	✓	✓✓
Privacy & Data Protection	✓✓	–	✓
Accountability Mechanisms	✓	✓	✓✓
Human Oversight	✓	✓✓	✓

Ratings reflect the degree to which each policy establishes explicit mechanisms for operationalizing each principle: ✓✓ strong (explicit mechanisms with enforcement), ✓ moderate (addressed but incomplete or internally enforced), – weak (acknowledged without substantive requirement)

Fairness is rated moderate across all three because each framework addresses only one dimension of bias while leaving others ungoverned — no framework earns strong here. Transparency is strong for Texas because Section 552.051 creates legally enforceable consumer disclosure backed by civil penalties — the only externally compelled transparency of the three. Anthropic and OSU earn moderate because their disclosure obligations are real but unenforced. Privacy is strong for Anthropic because ASL-2 and ASL-3 Security Standards represent the most technically developed privacy architecture of the three frameworks — though notably protecting model weights rather than users. Texas earns moderate for its biometric and social scoring provisions. OSU earns weak because it establishes no new privacy requirements whatsoever. Accountability is strong for Texas given Attorney General enforcement, civil penalties up to $200,000, and consumer complaint mechanisms. Anthropic and OSU earn moderate — their accountability structures are real and layered but fully internal to their respective institutions. Human Oversight is strong for OSU because instructor authority is absolute, pedagogically grounded, and free of the automation contradictions that undercut Anthropic's position. Anthropic earns moderate given robust procedural oversight mechanisms undermined by its own scaling trajectory. Texas earns moderate for establishing a legal floor through Section 552.052 without mandating affirmative oversight processes.

Strengths & Limitations:

Anthropic's Responsible Scaling Policy

This section evaluates the actionability, enforceability, and structural gaps within three AI governance frameworks: Anthropic's Responsible Scaling Policy, Ohio State University's AI Use Guidance, and Texas House Bill 149. Each reflects a distinct governance model: corporate self-regulation, institutional policy, and statutory mandate. Their differences illuminate trade-offs between flexibility, consistency, and accountability.

Anthropic's Responsible Scaling Policy (RSP) is notable for its proactive risk assessment architecture. Rather than reacting to harm after deployment, it establishes capability thresholds that trigger heightened safeguards before catastrophic risks materialize. This technical risk assessment approach enhances actionability by tying governance decisions directly to model capabilities. In practice, safety mechanisms scale in proportion to assessed risk levels through its AI Safety Level standards, creating an internally coherent structure where more powerful systems face stricter containment, evaluation, and oversight. This proportionality promotes efficient resource allocation while aligning safety investment with technological advancement.

The policy also demonstrates iterative adaptability. Its evolution from v1.0 to v2.2 signals responsiveness to emerging evidence and shifting threat models. Such adaptability strengthens long-term relevance, particularly in a field where static rules quickly decay. Transparency further bolsters credibility. Anthropic publicly discloses its framework, publishes capability reports in redacted form, and incorporates external expert consultation. While not fully transparent, the visible architecture signals procedural seriousness.

Finally, the governance structure is comparatively robust for a private entity. Defined roles including a Responsible Scaling Officer and board-level oversight bodies clarify decision-making authority. Whistleblower protections for employees introduce an internal accountability mechanism that may surface concerns before escalation.

Despite its technical sophistication, the RSP remains voluntary. There is no external enforcement mechanism compelling compliance; Anthropic retains unilateral authority to amend or withdraw the policy. This absence of binding oversight constrains enforceability and limits public recourse.

The policy's scope is also narrow. By focusing predominantly on catastrophic and existential risks, it may underemphasize everyday harms such as bias in hiring tools or discriminatory lending algorithms. This reflects stakeholder prioritization but leaves distributive and civil rights concerns comparatively underdeveloped.

Self-assessment presents another structural vulnerability. Internal teams evaluate their own systems, creating potential conflicts of interest even with external consultation. Moreover, the policy applies solely to Anthropic's models and does not establish industry-wide standards. Consumers and affected populations lack standing to challenge determinations or appeal outcomes, creating an accountability gap between corporate governance and public impact.

Ohio State University AI Use Guidance

The AI Use Guidance at Ohio State University reflects an institutional governance approach designed for pedagogical flexibility. Instructor discretion allows policies to align with disciplinary norms and assignment-specific learning objectives. This context-sensitive structure enhances actionability by embedding governance directly within classroom practice rather than imposing rigid prohibitions.

The guidance emphasizes educational goals over restriction. Rather than banning AI tools outright, it frames their use around whether they enhance or undermine learning outcomes. This pedagogical focus supports academic integrity while acknowledging technological reality. Equity awareness is also explicit. The guidance encourages instructors to consider disparities in student access to AI tools, foregrounding fairness concerns within policy design.

Practical examples across disciplines provide operational clarity. Concrete use cases in creative writing, sociology discussions, and quantitative problem-solving translate abstract principles into applied guidance. A citation framework further increases transparency by requiring students to acknowledge AI contributions, reinforcing academic norms of attribution.

However, flexibility produces inconsistency. Instructor discretion can result in uneven standards across courses and departments, generating confusion and potentially inequitable enforcement. Students navigating multiple classes may face conflicting expectations. Enforcement depends largely on instructor vigilance and detection capacity. Without technical safeguards, violations are difficult to identify reliably. The policy relies on compliance rather than systemic monitoring, limiting enforceability. Scope is another constraint. The guidance governs student academic use but does not comprehensively address institutional deployment of AI in admissions, advising, or administrative decision-making. Governance of those contexts remains fragmented. Additionally, student voice appears limited in policy formation. The structure operates primarily through top-down authority, potentially reducing perceived legitimacy among those most directly regulated.

Texas HB 149

Texas HB 149 represents a statutory model of AI governance. Its principal strength lies in legal enforceability. Civil penalties of up to $200,000 and enforcement authority vested in the Attorney General create tangible consequences for noncompliance. Unlike voluntary frameworks, this law imposes binding obligations. The statute's scope is comparatively comprehensive. It covers consumer-facing AI systems, government deployment, biometric data practices, and discriminatory applications. Specific prohibited uses, including manipulation, social scoring, and unauthorized biometric capture, reduce interpretive ambiguity and clarify red lines. This precision enhances actionability for regulated entities.

The inclusion of a regulatory sandbox balances innovation and oversight by permitting controlled experimentation with modified requirements. Consumer empowerment mechanisms, such as complaint procedures, disclosure mandates, and appeal rights, introduce formal accountability pathways unavailable in purely private frameworks. Nevertheless, the statute introduces its own constraints. Discrimination claims require proof of intent rather than recognizing disparate impact alone. This high evidentiary threshold may limit practical enforcement, even where harm occurs.

The 60-day cure period allows violators to remedy infractions before penalties attach. While designed to encourage compliance, it may permit ongoing harm during remediation. Enforcement capacity presents another concern. Technical AI investigations demand specialized expertise, and resource limitations within the Attorney General's office could constrain oversight effectiveness.

Legal definitions may also lag technological evolution. Statutory language risks obsolescence as AI capabilities shift, creating regulatory gaps. Finally, compliance burdens may disproportionately affect smaller developers lacking dedicated legal teams, potentially consolidating market power among larger firms.

Comparative Insights & Cross-Sector Analysis

From the Responsible Scaling Policy at Anthropic, to the guidance provided by instructors at Ohio State University, to Texas H.B. 149, one clear pattern becomes evident: the strength of a particular policy is going to be much more about who has the power to enforce that policy. These three documents cover a spectrum from voluntary commitment to legal compulsion. Where a policy is defined within that spectrum will determine who is protected by that policy, and how effectively. The Responsible Scaling Policy at Anthropic is the most technically specific about the details of the policy, yet it has the weakest external accountability. The company itself must evaluate its models, write reports about the safeguards it has implemented, and make a determination as to whether it proceeds with its AI systems. There is no external party that can force the organization to comply. In contrast, Texas H.B. 149 is far less specific about the technical details of the policy, but is far more coercive. The Attorney General has the power to investigate, assess, and impose significant civil penalties. For the Ohio State University, the consequences are more significant, though only within the school. This comparison makes clear that the ambition for governance is not necessarily the same as the power to govern.

A second pattern that emerges from reviewing these three documents is that each framework focuses on the types of risks that are immediately visible to the organization, while also ignoring others. Instead, the frameworks indicate three separate and mostly non-overlapping risk universes. Anthropic chose to focus on catastrophic, large-scale harms. For example, the thresholds established for CBern and autonomous AI are intended to prevent outcomes like the development of weapons of mass destruction and superintelligent AI. From the perspective of the individuals who coded and developed these systems, the concern for risks that could lead to the end of humanity or civilization is paramount. However, more routine types of harms, such as those related to hiring discrimination, biased policing, and manipulative recommendation systems, are overlooked. OSU's framework indicates that the primary concern is academic integrity. Students using tools such as ChatGPT in ways that are not disclosed to the instructor and outsourcing aspects of their education to AI tools are a major concern for the instructors at OSU. The framework provides guidelines for how to prevent plagiarism and encourage proper citation and attribution. However, there is no discussion of how the framework would govern the use of AI in more institutional contexts. Texas H.B. 149 focuses on the harms that result from consumer and civil rights being violated. For example, Texas is concerned with issues like manipulation of human behavior, mass surveillance, and discrimination. These are harms that are more documented and taking place in the present, rather than emerging frontier technologies. The result is that the frameworks exhibit some tunnel vision. Anthropic is very focused on long-term catastrophic outcomes but ignores more routine types of harm. The Ohio State University focuses on misuse in academic settings but ignores more institutional uses of AI. The state is focused on consumer rights and harms, but fails to consider the impact that AI might have on frontier technologies and development.

A third pattern that emerges from the reading of the Responsible Scaling Policy, the instructor guidance from Ohio State University, and Texas H.B. 149 is the use of the concept of "transparency" in different and distinct ways. The terms are used similarly, but their meaning and enforceability are different. Anthropic uses transparency as a means of gaining legitimacy for the company. The organization commits to publishing reports on the capabilities and safeguards of its models, maintaining a public changelog, having third-party reviews of its processes, and notifying government entities of the deployment of models with specific capabilities. The level of detail in these disclosures is unusually high for a private company. However, the company decides what gets published, when it gets published, and what gets redacted. The commitment can also be changed at any time by the company. OSU requires that students disclose to the instructors the use of tools such as ChatGPT and the way in which they used the tool to perform their work. This promotes the students' critical engagement with the AI. However, the requirement is limited to the students disclosing to the instructor; the instructors do not have to disclose their use of the tools. This level of transparency only occurs within the classroom. Finally, Texas H.B. 149 requires that entities such as government agencies and service providers make their use of AI transparent to the end users of their services. This involves publishing a notice of the use of AI at or before the user's interaction with the service, and the penalties for failing to do so. This regulation gives the Attorney General the power to enforce these requirements. The more durable form of transparency emerges from the analysis of these three case studies. The voluntary commitment to transparency is reliant on the continuing goodwill of the parties involved. The institutional commitment is contextual. The legally mandated commitment to transparency, though, is the one that will create lasting accountability.

At the core of each of these frameworks is a simple question: who has the power to define, investigate, and act upon AI violations? A comparison of the three frameworks reveals different answers to this critical question. The RSP gives all of its power to the organization itself. The leadership of the organization will define the violations, investigate them, and determine the action to be taken. End users have no power within this system. OSU places the power to define and investigate violations within the faculty members and internal review bodies. While students do have the ability to appeal to the faculty, they do not have any power to define or investigate. The most significant difference is in H.B. 149, which gives power to the consumer to file a complaint with the Attorney General. The Attorney General has the power to investigate and impose penalties on the violator. The type of governance that is used determines how power is distributed. The voluntary systems give all power to the developers, the institutional systems give it to the gatekeepers, and the legal frameworks give it to the public. These different systems determine who has a voice in AI governance.

In any given AI governance framework, there is a tension between innovation and protection. The following is a comparison of three of these frameworks: the Responsible Scaling Policy by Anthropic, the guidance from Ohio State University, and Texas House Bill 149. Anthropic's Responsible Scaling Policy is the most permissive on innovation. It allows for the continued development of models, provided the capabilities are not exceeded. Even if the thresholds are exceeded, additional safeguards are required rather than the development being stopped. However, there is little consideration of the potential harm that can result from these sub-catastrophic outcomes; the risk is placed upon the current users of the model. Ohio State University falls somewhere in the middle. The instructors will have to decide how the AI model can be used within the classroom setting. This provides some leeway for the innovators, although it places the onus on the students. The instructors have control, but the students have to comply. Texas House Bill 149 is the most protection-oriented. There are clear prohibitions on certain uses of AI models, but there is also a regulated sandbox in which experimentation is permitted. While innovation is allowed, there are clear limits on how it can be pursued. While this offers strong protection for the consumers, it increases the burden on the developers to comply with these regulations. In each of these frameworks, there are clearly trade-offs to be made. The framework indicates the priorities for the organization. While none of them offer a resolution of the tension between innovation and protection, the different structures place the balance of benefits and risks differently.

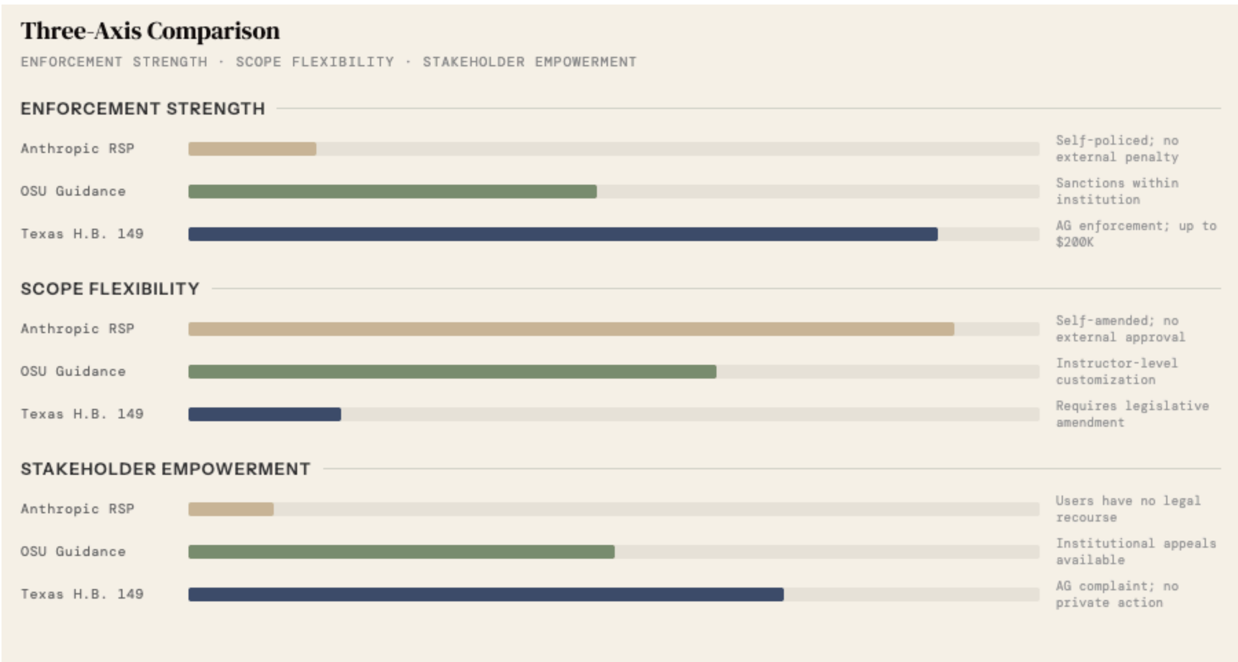

The comparison of the three frameworks reveals three main insights. The first insight is that the level of technical sophistication is inversely correlated with the degree of accountability. The most detailed and rigorous documentation of the requirements and purposes of Anthropic's RSP is matched with the weakest level of accountability; Texas H.B. 149, which says little about the technical operation of the AI system, is the only framework that can secure compliance from an unwilling actor. The second insight is that the blind spots of each framework coincide with the strengths of the others. The types of harm that are ignored by Anthropic are the types that Texas seeks to address; the institutional uses of AI that OSU does not see fall into the gap created by Texas's exemption of universities; the types of risk at the frontier that neither OSU nor Texas considers are the very category that Anthropic regulates. The third and final insight is that the level of transparency within each framework flows towards increasing institutional power rather than away from it. The summaries of the user reports are redacted before being given to the Board of Directors; the students at OSU are required to disclose the use of AI, but the instructors are not required to do so. What exists in these frameworks is not a system of governance but three separate and non-communicating frameworks, the gaps in each of which suggest that there was no process during their design that considered the full scope of these issues.

References

Anthropic. (2025, May 14). *Responsible scaling policy: Version 2.2*. https://www.anthropic.com/rsp-updates

Capriglione, G., et al. (2025). *Texas Responsible Artificial Intelligence Governance Act*, H.B. 149, 89th Texas Legislature, Regular Session. Texas Legislature Online. https://capitol.texas.gov/tlodocs/89R/analysis/html/HB00149S.htm

Ohio State University, College of Arts and Sciences Office of Distance Education. (n.d.). *Crafting policy for student use of artificial intelligence*. https://ascode.osu.edu/distance-learning-curriculum/crafting-policy-student-use-artificial-intelligence

AI Use Disclosure

All original analysis, comparative insights, and interpretive judgments in this paper reflect the independent work of the authoring team. Anthropic's Claude assisted with structural outlining, section organization, and identifying relevant provisions across policy documents for team review. QuillBot's grammar checker was applied to improve mechanical clarity and consistency across drafts.